

## # Zero-Trust Anomaly Detection System

### ## Project Report

**\*\*Project Title:\*\*** Zero-Trust Anomaly Detection in Authentication Logs  
**\*\*Date:\*\*** November 2024  
**\*\*Version:\*\*** 1.0

---

### ## Executive Summary

This report documents the development and implementation of a Zero-Trust Anomaly Detection System designed to identify suspicious authentication behaviors in real-time. The project successfully delivered a machine learning-based solution that detects anomalies such as impossible travel, off-hours logins, unusual resource access, and large data transfers using a combination of unsupervised learning models (Isolation Forest, One-Class SVM, and Autoencoder).

The system achieved a 59–63% accuracy rate across different models, with the Isolation Forest model demonstrating the best balance of precision and recall. The implementation includes a real-time REST API, interactive dashboard, and automated email alerting capabilities, providing a comprehensive security monitoring solution.

### **\*\*Key Achievements:\*\***

- Successfully implemented three ML models for anomaly detection
- Developed real-time prediction API with sub-second response times
- Created interactive dashboard for security analysts
- Integrated automated alerting system
- Achieved operational deployment readiness

---

### ## 1. Analysis of the Initial Problem

#### ### 1.1 Problem Statement

Organizations face increasing cybersecurity threats from sophisticated attackers who exploit authentication systems to

gain unauthorized access. Traditional security approaches rely on perimeter defenses and signature-based detection, which are insufficient against modern threats such as:

- **Credential Theft and Account Takeover:** Attackers use stolen credentials to access systems, appearing as legitimate users
- **Insider Threats:** Malicious or compromised insiders with valid credentials
- **Advanced Persistent Threats (APTs):** Long-term, stealthy attacks that evade traditional detection
- **Zero-Day Exploits:** Unknown attack patterns that bypass signature-based systems

### ### 1.2 Current State Analysis

#### **Existing Challenges:**

1. **Reactive Security Posture:** Traditional systems detect threats only after they occur, leading to delayed response times
2. **High False Positive Rates:** Rule-based systems generate numerous false alarms, causing alert fatigue
3. **Limited Behavioral Analysis:** Systems focus on known attack patterns rather than behavioral anomalies
4. **Manual Investigation Overhead:** Security teams spend significant time investigating false positives
5. **Lack of Real-Time Capabilities:** Batch processing delays threat detection and response

#### **Data Characteristics:**

- Dataset: 50,000 authentication events
- Anomaly Rate: ~40% (20,241 anomalies out of 50,000 events)
- Anomaly Types: 16 distinct categories including:
  - Impossible travel
  - Off-hours login
  - Multiple failed logins
  - Large data transfer
  - Unusual resource access
  - Combinations of the above

### ### 1.3 Business Impact

#### **Quantifiable Risks:**

- **Mean Time to Detection (MTTD):** Without automated detection, threats may go undetected for days or weeks
- **Data Breach Costs:** Average cost of a data breach exceeds \$4.45 million (2023 IBM Security Report)
- **Operational Disruption:** Security incidents cause downtime and productivity loss
- **Compliance Violations:** Failure to detect and respond to threats can result in regulatory penalties

### **Qualitative Risks:**

- Reputation damage from security breaches
- Loss of customer trust
- Intellectual property theft
- Competitive disadvantage

## **1.4 Solution Requirements**

The solution needed to address:

1. **Real-Time Detection:** Identify anomalies as they occur, not in retrospect
2. **Zero-Trust Architecture:** Treat all events as potentially suspicious until verified
3. **Explainability:** Provide clear reasoning for anomaly classifications
4. **Scalability:** Handle high-volume authentication events
5. **Integration:** Work with existing security infrastructure
6. **Usability:** Enable security analysts to quickly investigate and respond

---

## **2. Discussion of Improvement Opportunities**

### **2.1 Model Performance Enhancements**

#### **Current Performance:**

- Isolation Forest: 59% accuracy, 0.59 F1-score
- One-Class SVM: 59% accuracy, 0.59 F1-score
- Autoencoder: 63% accuracy, but low anomaly recall (10%)

#### **Improvement Opportunities:**

##### **1. Ensemble Methods**

- Combine predictions from multiple models using voting or stacking
- Expected improvement: 5–10% accuracy increase
- Implementation complexity: Medium

## 2. **\*\*Feature Engineering\*\***

- Add temporal features (time since last login, login frequency)
- Include user behavior baselines (average bytes transferred per user)
- Geographic features (distance from previous location)
- Expected improvement: 8–15% accuracy increase
- Implementation complexity: Low–Medium

## 3. **\*\*Hyperparameter Optimization\*\***

- Systematic grid search or Bayesian optimization
- Fine-tune contamination rates, kernel parameters, and neural network architecture
- Expected improvement: 3–7% accuracy increase
- Implementation complexity: Medium

## 4. **\*\*Advanced Deep Learning Models\*\***

- LSTM/GRU for sequence modeling of user behavior
- Transformer-based models for complex pattern recognition
- Expected improvement: 10–20% accuracy increase
- Implementation complexity: High

### ### 2.2 System Architecture Improvements

#### **\*\*Current Limitations:\*\***

- Kafka integration is optional and may not be available
- No persistent storage for historical predictions
- Limited model versioning and A/B testing capabilities

#### **\*\*Improvement Opportunities:\*\***

##### 1. **\*\*Database Integration\*\***

- Store all predictions and events in a time-series database (InfluxDB, TimescaleDB)
- Enable historical analysis and model retraining
- Implementation complexity: Medium

## 2. \*\*Model Versioning and Deployment\*\*

- Implement MLflow or similar for model registry
- Enable canary deployments and gradual rollouts
- A/B testing framework for model comparison
- Implementation complexity: Medium-High

## 3. \*\*Enhanced Streaming Architecture\*\*

- Guaranteed Kafka integration with error handling
- Event replay capabilities for model retraining
- Dead letter queue for failed predictions
- Implementation complexity: Medium

## 4. \*\*Microservices Architecture\*\*

- Separate model serving, preprocessing, and alerting services
- Independent scaling of components
- Better fault isolation
- Implementation complexity: High

### ### 2.3 Operational Improvements

#### \*\*Current Gaps:\*\*

- Manual model retraining process
- Limited monitoring and observability
- No automated model drift detection

#### \*\*Improvement Opportunities:\*\*

##### 1. \*\*Automated Model Retraining Pipeline\*\*

- Scheduled retraining on new data
- Automated feature drift detection
- Model performance monitoring
- Implementation complexity: Medium-High

##### 2. \*\*Comprehensive Monitoring\*\*

- Real-time model performance metrics
- Prediction latency tracking
- Alert volume and false positive rate monitoring
- Implementation complexity: Medium

### **3. \*\*Model Explainability Enhancement\*\***

- SHAP integration already present, but can be expanded
- LIME for local explanations
- Counterfactual explanations
- Implementation complexity: Low-Medium

### **4. \*\*Integration with Security Orchestration\*\***

- SOAR (Security Orchestration, Automation, and Response) integration
- Automated response actions (account lockout, IP blocking)
- SIEM integration for centralized logging
- Implementation complexity: High

## **### 2.4 User Experience Improvements**

### **\*\*Current State:\*\***

- Dashboard provides good visualization but could be enhanced
- Limited filtering and search capabilities
  - No bulk operations for analysts

### **\*\*Improvement Opportunities:\*\***

#### **1. \*\*Advanced Dashboard Features\*\***

- Real-time streaming updates without manual refresh
- Customizable alert rules and thresholds
- User-specific dashboards and saved filters
- Implementation complexity: Medium

#### **2. \*\*Investigation Workflow Tools\*\***

- Case management system for tracking investigations
- Collaboration features for security teams
- Integration with ticketing systems
- Implementation complexity: Medium-High

#### **3. \*\*Mobile Application\*\***

- Mobile alerts and basic dashboard access
- Push notifications for critical anomalies
- Implementation complexity: Medium

## **## 3. Business Case for Identified Improvements**

### ### 3.1 Financial Justification

#### **\*\*Cost-Benefit Analysis:\*\***

#### **\*\*Investment Required:\*\***

- Development resources: 3–6 months of engineering time
- Infrastructure: Additional compute for model training and serving (~\$500–2,000/month)
- Third-party tools: MLflow, monitoring tools (~\$200–500/month)
- **\*\*Total Estimated Investment: \$50,000 – \$150,000\*\***

#### **\*\*Expected Benefits:\*\***

##### **1. \*\*Reduced Security Incident Costs\*\***

- Current: Average detection time of 7–14 days
- Improved: Real-time detection reduces MTTD to minutes
- **\*\*Savings: \$200,000 – \$500,000 per prevented breach\*\***

##### **2. \*\*Operational Efficiency\*\***

- Reduced false positive investigation time: 20–30 hours/week saved
- Automated alerting reduces manual monitoring: 15–20 hours/week saved
- **\*\*Annual Savings: \$150,000 – \$250,000 in labor costs\*\***

##### **3. \*\*Compliance and Risk Mitigation\*\***

- Reduced risk of regulatory fines
- Improved audit trail and reporting
- **\*\*Value: \$50,000 – \$200,000 in avoided penalties\*\***

#### **\*\*ROI Calculation:\*\***

- **\*\*Total Annual Benefits: \$400,000 – \$950,000\*\***
- **\*\*Total Investment: \$50,000 – \$150,000\*\***
- **\*\*ROI: 267% – 1,800%\*\***
- **\*\*Payback Period: 1–3 months\*\***

### ### 3.2 Strategic Value

#### **\*\*Competitive Advantages:\*\***

1. **Proactive Security Posture:** Early threat detection provides competitive advantage
2. **Customer Trust:** Enhanced security builds customer confidence
3. **Innovation Leadership:** Demonstrates commitment to cutting-edge security practices
4. **Scalability:** System can grow with business needs

#### **Risk Mitigation:**

- Reduced exposure to cyber threats
- Better compliance with security regulations (GDPR, SOC 2, ISO 27001)
- Protection of intellectual property and sensitive data
- Business continuity assurance

### ### 3.3 Prioritization Framework

#### **High Priority (Quick Wins):**

1. Feature engineering improvements (Low complexity, high impact)
2. Hyperparameter optimization (Medium complexity, medium-high impact)
3. Enhanced monitoring and observability (Medium complexity, high operational value)

#### **Medium Priority (Strategic Investments):**

1. Automated model retraining pipeline (High complexity, high long-term value)
2. Database integration for historical analysis (Medium complexity, medium impact)
3. Advanced dashboard features (Medium complexity, high user value)

#### **Low Priority (Future Enhancements):**

1. Microservices architecture (High complexity, scalability benefit)
2. Mobile application (Medium complexity, convenience feature)
3. SOAR integration (High complexity, advanced automation)



## **## 4. Project Plan: Scope, Key Deliverables, and Suggested Changes**

### **### 4.1 Project Scope**

#### **\*\*In-Scope:\*\***

- Development and deployment of ML-based anomaly detection system
- Real-time API for prediction serving
- Interactive dashboard for security analysts
- Email alerting system
- Model training and evaluation framework
- Documentation and operational runbooks

#### **\*\*Out-of-Scope (Future Phases):\*\***

- Automated response actions (account lockout, IP blocking)
- Integration with external SIEM systems
- Mobile application development
- Multi-tenant architecture
- Advanced threat intelligence integration

### **### 4.2 Key Deliverables**

#### **\*\*Phase 1: Foundation (Completed)\*\***

- Data preprocessing pipeline
- Three ML models (Isolation Forest, One-Class SVM, Autoencoder)
- Model evaluation and comparison
- Basic REST API (FastAPI)
- Streamlit dashboard
- Email alerting system

#### **\*\*Phase 2: Enhancement (Recommended)\*\***

-  Feature engineering improvements
-  Hyperparameter optimization
-  Model ensemble implementation
-  Enhanced monitoring and logging

-  Database integration for historical data

### **\*\*Phase 3: Advanced Features (Future)\*\***

-  Automated model retraining pipeline
-  Advanced explainability features
-  SOAR integration
-  Performance optimization and scaling

## **#### 4.3 Suggested Changes and Improvements**

### **#### 4.3.1 Immediate Improvements (Next 1–2 Months)**

#### **\*\*1. Feature Engineering Enhancement\*\***

- **Action:** Add temporal and behavioral features
- **Deliverable:** Enhanced feature set with 15–20 features
- **Timeline:** 2–3 weeks
- **Resources:** 1 data scientist, 1 engineer

#### **\*\*2. Model Performance Optimization\*\***

- **Action:** Systematic hyperparameter tuning
- **Deliverable:** Optimized models with 5–10% accuracy improvement
- **Timeline:** 3–4 weeks
- **Resources:** 1 data scientist

#### **\*\*3. Monitoring and Observability\*\***

- **Action:** Implement comprehensive logging and metrics
- **Deliverable:** Dashboard with real-time performance metrics
- **Timeline:** 2–3 weeks
- **Resources:** 1 engineer

### **#### 4.3.2 Short-Term Improvements (3–6 Months)**

#### **\*\*1. Database Integration\*\***

- **Action:** Integrate time-series database for historical storage
- **Deliverable:** Persistent storage with query capabilities
- **Timeline:** 4–6 weeks
- **Resources:** 2 engineers

## **\*\*2. Automated Retraining Pipeline\*\***

- **Action:** Build CI/CD pipeline for model retraining
- **Deliverable:** Automated weekly/monthly retraining
- **Timeline:** 6–8 weeks
- **Resources:** 2 engineers, 1 data scientist

## **\*\*3. Enhanced Dashboard Features\*\***

- **Action:** Add real-time updates, advanced filtering, custom alerts
- **Deliverable:** Production-ready dashboard
- **Timeline:** 4–6 weeks
- **Resources:** 1 frontend engineer, 1 backend engineer

### **#### 4.3.3 Long-Term Improvements (6–12 Months)**

## **\*\*1. Advanced ML Models\*\***

- **Action:** Implement LSTM/Transformer models for sequence analysis
- **Deliverable:** Next-generation anomaly detection models
- **Timeline:** 8–12 weeks
- **Resources:** 2 data scientists, 1 ML engineer

## **\*\*2. SOAR Integration\*\***

- **Action:** Integrate with security orchestration platform
- **Deliverable:** Automated response capabilities
- **Timeline:** 10–12 weeks
- **Resources:** 2 engineers, 1 security specialist

## **\*\*3. Microservices Architecture\*\***

- **Action:** Refactor to microservices for better scalability
- **Deliverable:** Scalable, distributed system
- **Timeline:** 12–16 weeks
- **Resources:** 3–4 engineers

### **### 4.4 Risk Management**

#### **\*\*Technical Risks:\*\***

- **Model Performance Degradation:** Mitigation through continuous monitoring and automated retraining

- **Scalability Issues:** Mitigation through load testing and architecture improvements
- **Data Quality Issues:** Mitigation through data validation and quality checks

### **Operational Risks:**

- **Alert Fatigue:** Mitigation through intelligent threshold tuning and alert prioritization
- **System Downtime:** Mitigation through redundancy and failover mechanisms
- **Resource Constraints:** Mitigation through cloud auto-scaling

### **Business Risks:**

- **Changing Requirements:** Mitigation through agile development and regular stakeholder communication
- **Budget Constraints:** Mitigation through phased approach and ROI demonstration
- **Competing Priorities:** Mitigation through clear business case and executive sponsorship

---

## ## 5. Project Tracking and Success Metrics

### ### 5.1 Key Performance Indicators (KPIs)

#### #### 5.1.1 Model Performance Metrics

##### **Primary Metrics:**

- **Accuracy:** Target >70% (Current: 59–63%)
- **Precision:** Target >75% (Current: 50–81%)
- **Recall:** Target >70% (Current: 10–59%)
- **F1-Score:** Target >72% (Current: 18–59%)
- **ROC-AUC:** Target >0.80 (Current: ~0.60)

**Measurement Frequency:** Weekly during development, monthly in production

#### #### 5.1.2 System Performance Metrics

##### **Operational Metrics:**

- **API Response Time:** Target <100ms (P95)
- **System Uptime:** Target >99.9%
- **Throughput:** Target >1,000 requests/second
- **Error Rate:** Target <0.1%

**Measurement Frequency:** Real-time monitoring with daily reports

#### #### 5.1.3 Business Impact Metrics

##### **Security Metrics:**

- **Mean Time to Detect (MTTD):** Target <5 minutes
- **Mean Time to Respond (MTTR):** Target <30 minutes
- **False Positive Rate:** Target <10%
- **True Positive Rate:** Target >85%

##### **Operational Efficiency:**

- **Investigation Time per Alert:** Target <15 minutes (Current: ~45 minutes)
- **Alerts Processed per Analyst:** Target >50/day
- **Automated Response Rate:** Target >60% of low-risk anomalies

**Measurement Frequency:** Weekly reports, monthly trend analysis

#### ## 5.2 Tracking Methodology

##### **Data Collection:**

1. **Model Performance:** Automated evaluation on test sets and production data
2. **System Metrics:** Application Performance Monitoring (APM) tools
3. **Business Metrics:** Integration with ticketing and incident management systems
4. **User Feedback:** Regular surveys and interviews with security analysts

##### **Reporting:**

- **Daily:** System health and error rates
- **Weekly:** Model performance and business metrics
- **Monthly:** Comprehensive dashboard with trends and recommendations

- **\*\*Quarterly:\*\*** Executive summary with ROI analysis

### **### 5.3 Success Criteria**

#### **\*\*Phase 1 Success Criteria (Completed):\*\***

- Three ML models implemented and evaluated
- REST API operational with <200ms response time
- Dashboard functional with basic visualizations
- Email alerting system operational

#### **\*\*Phase 2 Success Criteria (Target):\*\***

- Model accuracy improved to >70%
- False positive rate reduced to <10%
- Automated retraining pipeline operational
- Database integration complete

#### **\*\*Phase 3 Success Criteria (Future):\*\***

- MTTD reduced to <5 minutes
- 85%+ true positive rate
- SOAR integration operational
- System handles 10,000+ events/second

### **### 5.4 Continuous Improvement Process**

#### **\*\*Feedback Loops:\*\***

1. **\*\*Model Performance Monitoring:\*\*** Automated alerts on performance degradation
2. **\*\*User Feedback Collection:\*\*** Regular surveys and feature requests
3. **\*\*Security Incident Analysis:\*\*** Post-incident reviews to identify detection gaps
4. **\*\*Competitive Analysis:\*\*** Monitoring industry best practices and new techniques

#### **\*\*Iteration Cycle:\*\***

- **\*\*Sprint Duration:\*\*** 2 weeks

- **Review Frequency:** End of each sprint
  - **Retrospective:** Monthly team retrospectives
  - **Model Retraining:** Weekly or monthly based on data volume
- 

## ## 6. Summary of Project Outcome and Lessons Learned

### ### 6.1 Project Outcomes

#### #### 6.1.1 Technical Achievements

##### **Successfully Delivered:**

1. **Multi-Model Anomaly Detection System:** Three different ML approaches providing diverse detection capabilities
2. **Real-Time Prediction API:** FastAPI-based service with sub-second response times
3. **Interactive Dashboard:** Streamlit application enabling security analysts to investigate anomalies
4. **Automated Alerting:** Email notifications for critical anomalies with detailed context
5. **Explainability Features:** SHAP integration providing model interpretability

##### **Performance Results:**

- Isolation Forest: 59% accuracy, balanced precision/recall
- One-Class SVM: 59% accuracy, similar performance to Isolation Forest
- Autoencoder: 63% accuracy but low anomaly recall (10%)
- Overall system capable of processing real-time authentication events

#### #### 6.1.2 Business Value Delivered

##### **Immediate Benefits:**

- Automated anomaly detection reducing manual monitoring effort
- Real-time threat detection capabilities
- Improved visibility into authentication patterns
- Foundation for advanced security analytics

##### **Strategic Value:**

- Zero-trust security architecture implementation

- Scalable platform for future enhancements
- Data-driven security decision making
- Enhanced security posture

## ### 6.2 Key Lessons Learned

### #### 6.2.1 Technical Lessons

#### \*\*1. Model Selection and Evaluation\*\*

- **Lesson:** Different models excel at different anomaly types
- **Insight:** Ensemble approaches may provide better overall performance
- **Application:** Consider model stacking or voting for production

#### \*\*2. Feature Engineering Importance\*\*

- **Lesson:** Simple features (hour, bytes\_transferred) are highly effective
- **Insight:** Temporal and behavioral features significantly improve detection
- **Application:** Invest in feature engineering before complex model architectures

#### \*\*3. Explainability is Critical\*\*

- **Lesson:** Security teams need to understand why an event is flagged
- **Insight:** SHAP values provide actionable insights for investigations
- **Application:** Prioritize explainability features in production systems

#### \*\*4. Real-Time Processing Challenges\*\*

- **Lesson:** Categorical encoding for new values requires careful handling
- **Insight:** Preprocessing pipeline must handle unseen categories gracefully
- **Application:** Implement robust data validation and fallback mechanisms

### #### 6.2.2 Process Lessons

#### \*\*1. Iterative Development Approach\*\*

- **Lesson:** Starting with multiple models provided valuable comparisons
- **Insight:** Rapid prototyping enabled quick learning and iteration
- **Application:** Continue agile approach with regular model updates

## **\*\*2. User-Centric Design\*\***

- **Lesson:** Dashboard usability directly impacts analyst productivity
- **Insight:** Security analysts need fast access to relevant information
- **Application:** Regular user feedback sessions essential for UX improvements

## **\*\*3. Operational Considerations\*\***

- **Lesson:** Email alerting requires careful configuration and testing
- **Insight:** Alert fatigue is a real concern with high false positive rates
- **Application:** Implement intelligent alert prioritization and threshold tuning

## **\*\*4. Documentation and Maintenance\*\***

- **Lesson:** Well-documented code and processes enable faster onboarding
- **Insight:** Model retraining procedures need to be clearly documented
- **Application:** Maintain comprehensive documentation and runbooks

## **#### 6.2.3 Business Lessons**

### **\*\*1. ROI Demonstration\*\***

- **Lesson:** Quantifiable metrics are essential for stakeholder buy-in
- **Insight:** Early wins (automated detection) provide immediate value
- **Application:** Track and report business metrics regularly

### **\*\*2. Phased Approach\*\***

- **Lesson:** Starting with MVP and iterating is more effective than big-bang delivery
- **Insight:** Each phase delivers incremental value
- **Application:** Continue phased enhancement approach

### **3. Integration Challenges**

- **Lesson:** Kafka integration optionality added complexity
- **Insight:** External dependencies should be clearly defined and tested
- **Application:** Minimize optional dependencies or provide clear alternatives

## **6.3 Challenges Encountered**

### **Technical Challenges:**

1. **Low Anomaly Recall in Autoencoder:** Model struggled to identify anomalies despite good overall accuracy
  - **Resolution:** Focused on Isolation Forest for production, continued Autoencoder research
2. **Categorical Encoding for New Values:** Handling unseen categories in real-time predictions
  - **Resolution:** Implemented dynamic encoding with fallback mechanisms
3. **Model Performance Optimization:** Balancing precision and recall
  - **Resolution:** Used F1-score as primary metric, tuned contamination rates

### **Operational Challenges:**

1. **Email Configuration Complexity:** Gmail App Password setup required multiple iterations
  - **Resolution:** Created detailed documentation and setup guides
2. **Dashboard Performance:** Large datasets caused slow rendering
  - **Resolution:** Implemented filtering and pagination
3. **Real-Time Data Synchronization:** Ensuring dashboard reflects latest events
  - **Resolution:** File-based approach with manual refresh (future: real-time streaming)

## **6.4 Recommendations for Future Work**

### **\*\*Immediate Priorities (Next Quarter):\*\***

1. **Feature Engineering:** Add temporal and behavioral features to improve model accuracy
2. **Hyperparameter Optimization:** Systematic tuning to achieve >70% accuracy
3. **Monitoring Enhancement:** Comprehensive observability for production operations

### **\*\*Short-Term Priorities (6 Months):\*\***

1. **Automated Retraining:** Implement CI/CD pipeline for model updates
2. **Database Integration:** Historical storage and analysis capabilities
3. **Advanced Dashboard:** Real-time updates and enhanced filtering

### **\*\*Long-Term Vision (12+ Months):\*\***

1. **Advanced ML Models:** LSTM/Transformer for sequence analysis
2. **SOAR Integration:** Automated response and orchestration
3. **Enterprise Scalability:** Microservices architecture for high-volume deployments

### **### 6.5 Conclusion**

The Zero-Trust Anomaly Detection System project successfully delivered a functional ML-based security monitoring solution. While model performance (59–63% accuracy) has room for improvement, the system provides a solid foundation for real-time anomaly detection with clear paths for enhancement.

The project demonstrated the value of:

- **Iterative Development:** Rapid prototyping and continuous improvement
- **Multi-Model Approach:** Diversity in detection methods
- **Explainability:** SHAP integration for actionable insights
- **User-Centric Design:** Dashboard and alerting tailored to security analysts

### **\*\*Key Success Factors:\*\***

- Clear problem definition and requirements
- Agile development methodology
- Focus on operational usability
- Comprehensive documentation

### **\*\*Next Steps:\*\***

1. Implement Phase 2 improvements (feature engineering, optimization)
2. Deploy to production with monitoring
3. Gather user feedback and iterate
4. Plan Phase 3 advanced features

The project has established a strong foundation for zero-trust security monitoring, with clear opportunities for enhancement and significant potential for business value through improved threat detection and operational efficiency.



### **## Appendices**

#### **### Appendix A: Technical Architecture**

### **\*\*Components:\*\***

- **Model Training:** Jupyter notebook with scikit-learn and TensorFlow
- **API Service:** FastAPI with joblib model loading
- **Dashboard:** Streamlit with Plotly visualizations
- **Alerting:** SMTP email integration
- **Streaming:** Optional Kafka integration

### **\*\*Technology Stack:\*\***

- Python 3.11+
- scikit-learn, TensorFlow, pandas, numpy
- FastAPI, Streamlit
- Kafka (optional), Docker

#### **### Appendix B: Model Performance Summary**

Model	Accuracy	Precision	Recall	F1-Score	
ROC-AUC					
	-----	-----	-----	-----	-----
	-----	-----	-----	-----	-----

Isolation Forest	59%	50–68%	59%	54–63%	
~0.60	[REDACTED]				
One-Class SVM	59%	50–68%	59%	54–59%	
~0.60	[REDACTED]				
Autoencoder	63%	62–81%	10–98%	18–76%	
~0.54	[REDACTED]				

### ### Appendix C: Project Timeline

#### **\*\*Phase 1 (Completed):\*\* 8–10 weeks**

- Weeks 1–2: Data exploration and preprocessing
- Weeks 3–5: Model development and evaluation
- Weeks 6–7: API and dashboard development
- Weeks 8–10: Integration, testing, and documentation

#### **\*\*Phase 2 (Planned):\*\* 12–16 weeks**

- Feature engineering and optimization
- Database integration
- Enhanced monitoring
- Automated retraining pipeline

---

**\*\*Document End\*\***